

Problem Statement

Business Context

Workplace safety in hazardous environments like construction sites and industrial plants is crucial to prevent accidents and injuries. One of the most important safety measures is ensuring workers wear safety helmets, which protect against head injuries from falling objects and machinery. Non-compliance with helmet regulations increases the risk of serious injuries or fatalities, making effective monitoring essential, especially in large-scale operations where manual oversight is prone to errors and inefficiency.

To overcome these challenges, SafeGuard Corp plans to develop an automated image analysis system to detect whether workers are wearing safety helmets. This system will improve safety enforcement, ensuring compliance and reducing the risk of head injuries. By automating helmet monitoring, SafeGuard aims to enhance efficiency, scalability, and accuracy, ultimately fostering a safer work environment while minimizing human error in safety oversight.

Objective

As a data scientist at SafeGuard Corp, you are tasked with developing an image classification model that classifies images into one of two categories:

- **With Helmet:** Workers wearing safety helmets.
- **Without Helmet:** Workers not wearing safety helmets.

The ultimate objective is to prepare the model for deployment as part of an automated monitoring system. This system will enable real-time analysis of images and assist in ensuring compliance with workplace safety regulations.

Data Description

The dataset consists of **631 images**, equally divided into two categories:

- **With Helmet:** 311 images showing workers wearing helmets.
- **Without Helmet:** 320 images showing workers not wearing helmets.

Dataset Characteristics:

- **Variations in Conditions:** Images include diverse environments such as construction sites, factories, and industrial settings, with variations in lighting, angles, and worker postures to simulate real-world conditions.
- **Worker Activities:** Workers are depicted in different actions such as standing, using tools, or moving, ensuring robust model learning for various scenarios.

Installing and Importing the Necessary Libraries

```
!pip install tensorflow -q
```

```
import tensorflow as tf
print("Num GPUs Available:",
len(tf.config.list_physical_devices('GPU')))
print(tf.__version__)
```

```
Num GPUs Available: 1
2.18.0
```

```
import os
import random
import numpy as np
# Importing numpy for Matrix Operations
import pandas as pd
import seaborn as sns
import matplotlib.image as mpimg
# Importing pandas to read CSV files
import matplotlib.pyplot as plt
# Importing matplotlib for Plotting and visualizing images
import math
# Importing math module to perform mathematical operations
import cv2

# Tensorflow modules
import keras
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
# Importing the ImageDataGenerator for data augmentation
from tensorflow.keras.models import Sequential
# Importing the sequential module to define a sequential model
from tensorflow.keras.layers import
Dense,Dropout,Flatten,Conv2D,MaxPooling2D,BatchNormalization #
Defining all the layers to build our CNN Model
from tensorflow.keras.optimizers import Adam,SGD
# Importing the optimizers which can be used in our model
from sklearn import preprocessing
# Importing the preprocessing module to preprocess the data
from sklearn.model_selection import train_test_split
# Importing train_test_split function to split the data into train and test
from sklearn.metrics import confusion_matrix
# Importing confusion_matrix to plot the confusion matrix
from tensorflow.keras.models import Model
from keras.applications.vgg16 import VGG16
```

```

from tensorflow.keras.layers import SpatialDropout2D,
GlobalAveragePooling2D
from tensorflow.keras.utils import clear_session as cls

# Display images using OpenCV
from google.colab.patches import cv2_imshow

#Imports functions for evaluating the performance of machine learning
models
from sklearn.metrics import confusion_matrix, f1_score, accuracy_score,
recall_score, precision_score, classification_report
from sklearn.metrics import mean_squared_error as mse
# Importing cv2_imshow from google.patches to display images

# Ignore warnings
import warnings
warnings.filterwarnings('ignore')

# Set the seed using keras.utils.set_random_seed. This will set:
# 1) `numpy` seed
# 2) backend random seed
# 3) `python` random seed
tf.keras.utils.set_random_seed(812)

```

Data Overview

##Loading the data

```

from google.colab import drive
drive.mount('/content/drive')
pathl='/content/drive/MyDrive/Dataset/Labels_proj.csv'
pathim='/content/drive/MyDrive/Dataset/images_proj.npy'

Mounted at /content/drive

labels=pd.read_csv(pathl)
labels.head()

{"summary":{"\n  \"name\": \"labels\", \n  \"rows\": 631, \n
  \"fields\": [\n    {\n      \"column\": \"Label\", \n
  \"properties\": {\n      \"dtype\": \"number\", \n      \"std\":
0, \n      \"min\": 0, \n      \"max\": 1, \n
  \"num_unique_values\": 2, \n      \"samples\": [\n      0, \n
1\n    ], \n      \"semantic_type\": \"\", \n
  \"description\": \"\" \n    } \n  ] \n
n} \", \"type\": \"dataframe\", \"variable_name\": \"labels\"}

```

```
images=np.load(pathim)
images.shape

(631, 200, 200, 3)
```

Exploratory Data Analysis


###Plotting random images from each of the classes and printing their corresponding labels.

```
la1=labels[labels['Label']==1].index.tolist()
la0=labels[labels['Label']==0].index.tolist()
lam=la1[:4]+la0[:4]
lam

[0, 1, 2, 3, 271, 272, 273, 274]

fig,ax=plt.subplots(2,4,figsize=(15,15))

ax=ax.flatten()
for i,j in zip(ax,lam):
    i.imshow(images[j])
    i.axis('off')
    if j in la1:
        i.set_title('With Helmet')
    else:
        i.set_title('Without Helmet')
plt.tight_layout()
plt.show()
```



Observations


- From the images, we observe a blue tint, which occurs because matplotlib prints in RGB format, while the image might be in BGR format.

Checking for class imbalance

```
a=sns.countplot(data=labels,x='Label')
for i in a.patches:

    a.text(
        x=(i.get_x() + i.get_width() / 2),
        y=i.get_height() + 5,
        s=f"{i.get_height()}",
        ha='center',
        va='center',
        size=12
    )

# Show the plot
plt.show()
```



Observations:

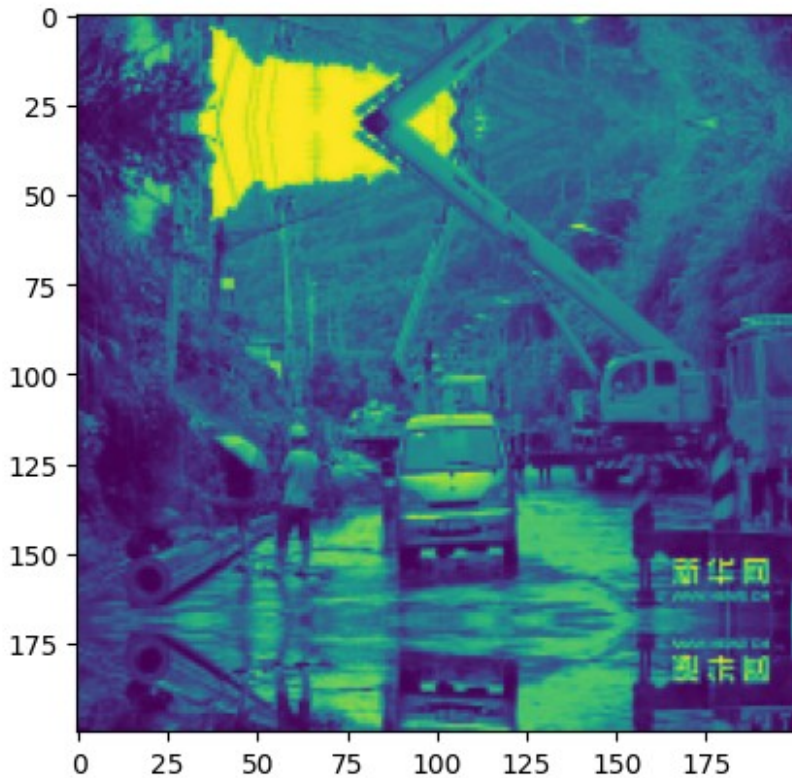
- The Class imbalance is not that very big to affect the performance of the model.

Data Preprocessing

Converting images to grayscale

```
imagegs=np.array([cv2.cvtColor(image,cv2.COLOR_BGR2GRAY) for image in
images])
imagegs=imagegs.reshape(631,200,200,1)
imagegs.shape

(631, 200, 200, 1)
plt.imshow(imagegs[0])
<matplotlib.image.AxesImage at 0x7c9699903950>
```



Splitting the dataset

```
x=images
y=labels['Label']

x_train,x_temp,y_train,y_temp=train_test_split(x,y,test_size=0.4,stratify=y,random_state=812)

x_val,x_test,y_val,y_test=train_test_split(x_temp,y_temp,test_size=0.5,
,stratify=y_temp,random_state=812)

x_val.shape,x_test.shape,x_train.shape

((126, 200, 200, 1), (127, 200, 200, 1), (378, 200, 200, 1))

y_val.value_counts(),y_test.value_counts(),y_train.value_counts()

(Label
0    64
1    62
Name: count, dtype: int64,
Label
0    64
1    63
Name: count, dtype: int64,
Label
0   192
```

```
1      186
Name: count, dtype: int64)
```

Data Normalization

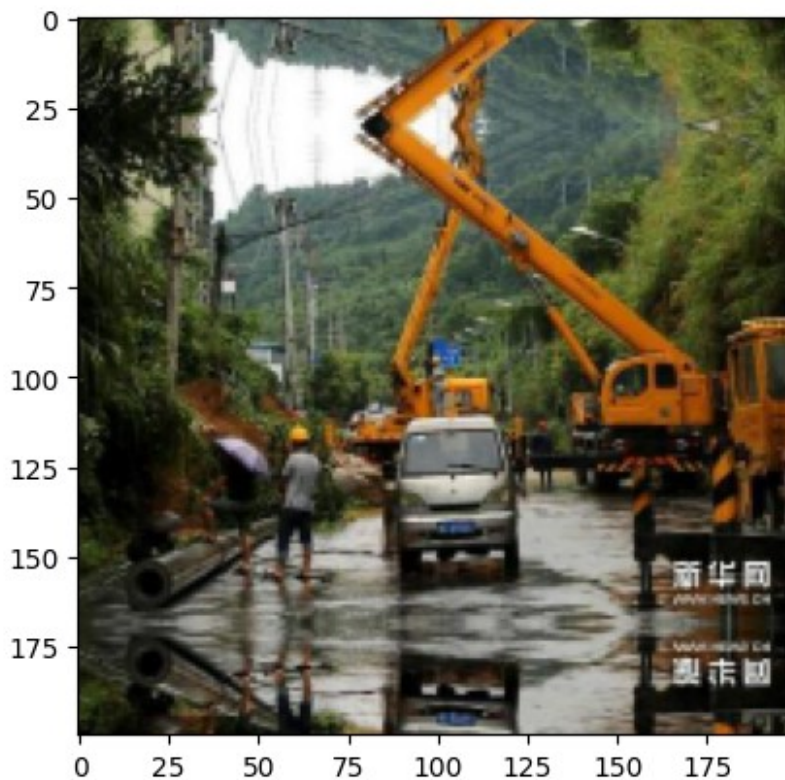
```
x_train_normalized = x_train.astype('float32')/255.0
x_val_normalized = x_val.astype('float32')/255.0
x_test_normalized = x_test.astype('float32')/255.0

y_train[0]
np.int64(1)
```

Data splitting & Normalisation for color data:

- Since VGG16 model works with three colour channels.

```
imagcol=np.array([cv2.cvtColor(image,cv2.COLOR_BGR2RGB) for image in
images])
imagcol.shape
(631, 200, 200, 3)
plt.imshow(imagcol[0])
<matplotlib.image.AxesImage at 0x7c96999df110>
```




```

xc=imagcol
yc=labels['Label']

xc_train,xc_temp,yc_train,yc_temp=train_test_split(xc,yc,test_size=0.4
,stratify=y,random_state=812)

xc_val,xc_test,yc_val,yc_test=train_test_split(xc_temp,yc_temp,test_si
ze=0.5,stratify=yc_temp,random_state=812)

xc_val.shape,xc_test.shape,xc_train.shape
((126, 200, 200, 3), (127, 200, 200, 3), (378, 200, 200, 3))

yc_val.value_counts(),yc_test.value_counts(),yc_train.value_counts()
(Label
0      64
1      62
Name: count, dtype: int64,
Label
0      64
1      63
Name: count, dtype: int64,
Label
0      192
1      186
Name: count, dtype: int64)

xc_train_normalized = xc_train.astype('float32')/255.0
xc_val_normalized = xc_val.astype('float32')/255.0
xc_test_normalized = xc_test.astype('float32')/255.0

```

Model Building

##Model Evaluation Criterion

The primary metric for evaluating the model will be **Recall** for the "Without Helmet" category.

Reason for Choosing Recall:

- **Critical Safety Focus:** Missing instances of workers not wearing helmets can result in severe safety risks. Prioritizing recall ensures that most non-compliant cases are detected, minimizing the risk of the workers getting injured.
- **Error Consequences:** False negatives (failing to detect a worker without a helmet) are more critical than false positives (incorrectly flagging a compliant worker), as undetected non-compliance could lead to accidents which could be even life threatening and also lead to legal issues.

Secondary Metrics:

- **Precision for "Without Helmet":** To avoid excessive false positives that could undermine system reliability.
- **F1 Score:** To balance recall and precision, ensuring an overall robust performance.
- **Accuracy:** For general performance evaluation.

Utility Functions

```
# defining a function to compute different metrics to check  
performance of a classification model built using statsmodels  
def model_performance_classification(model, predictors, target):  
    """  
        Function to compute different metrics to check classification  
        model performance  
  
        model: classifier  
        predictors: independent variables  
        target: dependent variable  
    """  
  
    # checking which probabilities are greater than threshold  
    pred = model.predict(predictors).reshape(-1)>0.5  
  
    target = target.to_numpy().reshape(-1)  
  
    acc = accuracy_score(target, pred) # to compute Accuracy  
    recall = recall_score(target, pred, average='weighted') # to  
compute Recall  
    precision = precision_score(target, pred, average='weighted') #  
to compute Precision  
    f1 = f1_score(target, pred, average='weighted') # to compute F1-  
score  
  
    # creating a dataframe of metrics  
    df_perf = pd.DataFrame({"Accuracy": acc, "Recall": recall,  
"Precision": precision, "F1 Score": f1}, index=[0],)  
  
    return df_perf  
  
def plot_confusion_matrix(model, predictors, target, ml=False):  
    """  
        Function to plot the confusion matrix  
  
        model: classifier  
        predictors: independent variables  
        target: dependent variable  
        ml: To specify if the model used is an sklearn ML model or not  
(True means ML model)  
    """
```

```

# checking which probabilities are greater than threshold
pred = model.predict(predictors).reshape(-1)>0.5

target = target.to_numpy().reshape(-1)

# Plotting the Confusion Matrix using confusion_matrix() function
which is also predefined tensorflow module
confusion_matrix = tf.math.confusion_matrix(target,pred)
f, ax = plt.subplots(figsize=(10, 8))
sns.heatmap(
    confusion_matrix,
    annot=True,
    linewidths=.4,
    fmt="d",
    square=True,
    xticklabels=["Without Helmet", "With Helmet"],
    yticklabels=["Without Helmet", "With Helmet"],
    ax=ax
)
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix")
plt.show()

# model for plotting loss
def plot_history(history):
    plt.figure(figsize=(15,5))
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
    plt.title('Model loss')
    plt.ylabel('Loss')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper right')
    plt.show()

```

##Model 1: Simple Convolutional Neural Network (CNN)

```

cls()

m1=Sequential()
m1.add(Conv2D(64,
(3,3),activation='relu',input_shape=(200,200,1),padding = 'same'))
m1.add(BatchNormalization())
m1.add(MaxPooling2D(2,2))

m1.add(Conv2D(32,(3,3),activation='relu',padding='same'))
m1.add(BatchNormalization())
m1.add(MaxPooling2D(2,2))

```

```

m1.add(Conv2D(32,(3,3),activation='relu',padding='same'))
m1.add(MaxPooling2D(2,2))
m1.add(Flatten())
m1.add(Dense(128,activation='relu'))
m1.add(BatchNormalization())
m1.add(Dropout(0.5))
m1.add(Dense(64,activation='relu'))
m1.add(BatchNormalization())
m1.add(Dropout(0.5))
m1.add(Dense(1,activation='sigmoid'))

m1.summary()

```

Model: "sequential"

Layer (type) Param #	Output Shape	
conv2d (Conv2D) 640	(None, 200, 200, 64)	
batch_normalization 256 (BatchNormalization)	(None, 200, 200, 64)	
max_pooling2d (MaxPooling2D) 0	(None, 100, 100, 64)	
conv2d_1 (Conv2D) 18,464	(None, 100, 100, 32)	
batch_normalization_1 128 (BatchNormalization)	(None, 100, 100, 32)	
max_pooling2d_1 (MaxPooling2D) 0	(None, 50, 50, 32)	
conv2d_2 (Conv2D)	(None, 50, 50, 32)	

9,248				
		max_pooling2d_2 (MaxPooling2D)	(None, 25, 25, 32)	
0				
		flatten (Flatten)	(None, 20000)	
0				
		dense (Dense)	(None, 128)	
2,560,128				
		batch_normalization_2	(None, 128)	
512		(BatchNormalization)		
		dropout (Dropout)	(None, 128)	
0				
		dense_1 (Dense)	(None, 64)	
8,256				
		batch_normalization_3	(None, 64)	
256		(BatchNormalization)		
		dropout_1 (Dropout)	(None, 64)	
0				
		dense_2 (Dense)	(None, 1)	
65				

Total params: 2,597,953 (9.91 MB)

Trainable params: 2,597,377 (9.91 MB)

Non-trainable params: 576 (2.25 KB)

```
optil=keras.optimizers.Adam(learning_rate=0.001)
m1.compile(optimizer=optil,loss='binary_crossentropy',metrics=['accuracy'])
```

```
h1=m1.fit(x_train_normalized,y_train,epochs=32,validation_data=(x_val_normalized,y_val),batch_size=32)
```

Epoch 1/32

12/12 _____ 18s 671ms/step - accuracy: 0.7948 - loss: 0.4489 - val_accuracy: 0.5317 - val_loss: 0.6626

Epoch 2/32

12/12 _____ 6s 54ms/step - accuracy: 0.9554 - loss: 0.1190 - val_accuracy: 0.5079 - val_loss: 0.9911

Epoch 3/32

12/12 _____ 1s 50ms/step - accuracy: 0.9581 - loss: 0.0884 - val_accuracy: 0.5079 - val_loss: 1.3227

Epoch 4/32

12/12 _____ 1s 50ms/step - accuracy: 0.9749 - loss: 0.0517 - val_accuracy: 0.5079 - val_loss: 1.5227

Epoch 5/32

12/12 _____ 1s 50ms/step - accuracy: 0.9913 - loss: 0.0314 - val_accuracy: 0.5079 - val_loss: 1.8202

Epoch 6/32

12/12 _____ 1s 50ms/step - accuracy: 0.9923 - loss: 0.0288 - val_accuracy: 0.5079 - val_loss: 1.8100

Epoch 7/32

12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss: 0.0158 - val_accuracy: 0.5079 - val_loss: 2.1666

Epoch 8/32

12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss: 0.0165 - val_accuracy: 0.5079 - val_loss: 2.4214

Epoch 9/32

12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss: 0.0114 - val_accuracy: 0.5079 - val_loss: 2.6409

Epoch 10/32

12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss: 0.0103 - val_accuracy: 0.5079 - val_loss: 2.7557

Epoch 11/32

12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss: 0.0087 - val_accuracy: 0.5079 - val_loss: 2.9510

Epoch 12/32

12/12 _____ 1s 49ms/step - accuracy: 1.0000 - loss: 0.0091 - val_accuracy: 0.5079 - val_loss: 2.9800

Epoch 13/32

12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss: 0.0054 - val_accuracy: 0.5079 - val_loss: 3.0975

Epoch 14/32

12/12 _____ 1s 60ms/step - accuracy: 1.0000 - loss: 0.0069 - val_accuracy: 0.5079 - val_loss: 3.2048

Epoch 15/32

```
12/12 _____ 1s 80ms/step - accuracy: 1.0000 - loss:
0.0053 - val_accuracy: 0.5079 - val_loss: 3.2756
Epoch 16/32
12/12 _____ 1s 102ms/step - accuracy: 1.0000 - loss:
0.0049 - val_accuracy: 0.5079 - val_loss: 3.2095
Epoch 17/32
12/12 _____ 2s 53ms/step - accuracy: 1.0000 - loss:
0.0049 - val_accuracy: 0.5079 - val_loss: 3.3308
Epoch 18/32
12/12 _____ 1s 50ms/step - accuracy: 0.9967 - loss:
0.0062 - val_accuracy: 0.5079 - val_loss: 3.2886
Epoch 19/32
12/12 _____ 1s 49ms/step - accuracy: 1.0000 - loss:
0.0055 - val_accuracy: 0.5079 - val_loss: 3.0630
Epoch 20/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0040 - val_accuracy: 0.5079 - val_loss: 3.0490
Epoch 21/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0027 - val_accuracy: 0.5079 - val_loss: 3.2925
Epoch 22/32
12/12 _____ 1s 51ms/step - accuracy: 1.0000 - loss:
0.0029 - val_accuracy: 0.5079 - val_loss: 3.3412
Epoch 23/32
12/12 _____ 1s 51ms/step - accuracy: 0.9891 - loss:
0.0125 - val_accuracy: 0.5079 - val_loss: 3.5700
Epoch 24/32
12/12 _____ 1s 49ms/step - accuracy: 1.0000 - loss:
0.0021 - val_accuracy: 0.5079 - val_loss: 3.5227
Epoch 25/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0029 - val_accuracy: 0.5079 - val_loss: 3.3244
Epoch 26/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0030 - val_accuracy: 0.5079 - val_loss: 3.0013
Epoch 27/32
12/12 _____ 1s 51ms/step - accuracy: 1.0000 - loss:
0.0018 - val_accuracy: 0.5159 - val_loss: 2.8879
Epoch 28/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0023 - val_accuracy: 0.5238 - val_loss: 2.9027
Epoch 29/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0021 - val_accuracy: 0.5317 - val_loss: 2.7379
Epoch 30/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
0.0020 - val_accuracy: 0.5635 - val_loss: 2.3283
Epoch 31/32
12/12 _____ 1s 50ms/step - accuracy: 1.0000 - loss:
```

0.0016 - val_accuracy: 0.6032 - val_loss: 1.9654

Epoch 32/32

12/12 ————— 1s 60ms/step - accuracy: 1.0000 - loss:

0.0013 - val_accuracy: 0.6111 - val_loss: 1.8073

m1tr=model_performance_classification(m1,x_train_normalized,y_train)


m1tr

12/12 ————— 2s 111ms/step

```
{"summary":{"\n  \"name\": \"m1tr\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.5952380952380952,\n        \"max\": 0.5952380952380952,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.5952380952380952\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.5952380952380952,\n        \"max\": 0.5952380952380952,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.5952380952380952\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.7747412008281573,\n        \"max\": 0.7747412008281573,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.7747412008281573\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"F1 Score\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.5115102748836582,\n        \"max\": 0.5115102748836582,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.5115102748836582\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  }},\n  \"type\": \"dataframe\",\n  \"variable_name\": \"m1tr\"}
```

plot_confusion_matrix(m1,x_train_normalized,y_train)

12/12 ————— 0s 10ms/step



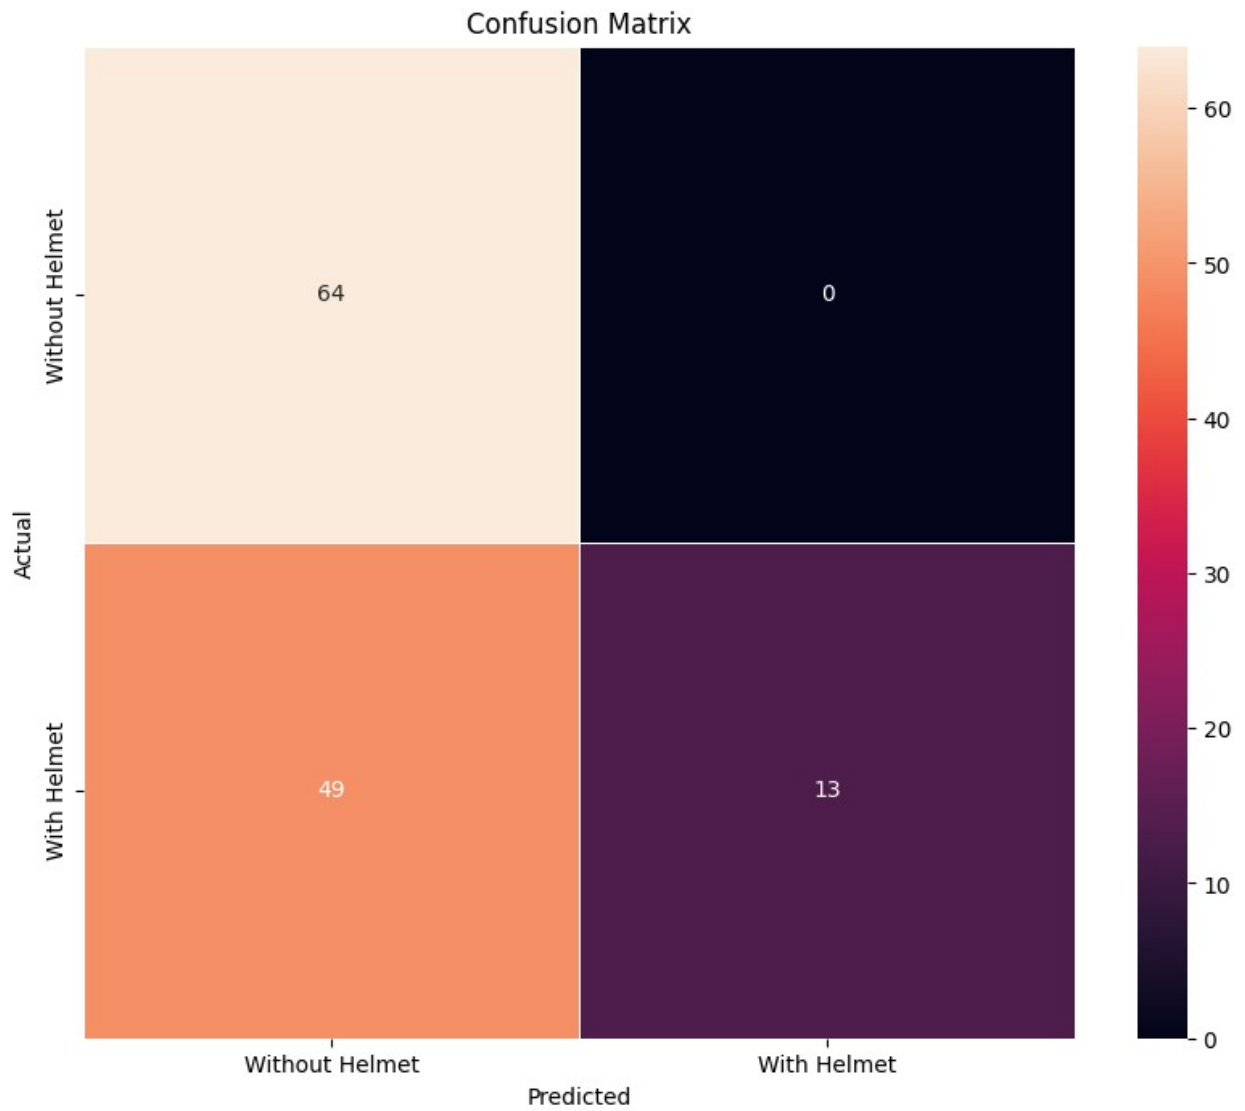
```
m1vl=model_performance_classification(m1,x_val_normalized,y_val)
m1vl
```

4/4 ————— 0s 92ms/step


```
{"summary":{"\n  \"name\": \"m1vl\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.6111111111111112,\n        \"max\": 0.6111111111111112,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.6111111111111112\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 0.6111111111111112,\n        \"max\": 0.6111111111111112,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          0.6111111111111112\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ]\n}}
```

```
plot_confusion_matrix(m1,x_val_normalized,y_val)
```

4/4 ————— 0s 15ms/step



```
plot_history(h1)
```



Vizualizing the predictions

```
pre = m1.predict(x_train_normalized)
f,ax=plt.subplots(4,2,figsize=(15,15))
for i,j in zip(ax.flatten(),range(8)):
    i.imshow(x_train[j])
    i.axis('off')
    if pre[j]>0.5:
        i.set_title(f'With Helmet:{pre[j]}')
    else:
        i.set_title(f'Without Helmet:{pre[j]}')
```

12/12 ————— 0s 11ms/step


Without Helmet:[0.01579482]




Without Helmet:[0.00044471]



Without Helmet:[0.00012029]




Without Helmet:[0.00234477]



Without Helmet:[0.00321866]



Without Helmet:[0.04541063]



Without Helmet:[0.00116103]



Without Helmet:[8.346523e-05]



Observations on model 1:

- As we observe the metrics evaluations and confusion matrix we can clearly see that recall is very low and metrics are not upto mark.
- The FN are fairly high with low recalls which means the model is not able to distinguish people With Helmets as FP are zero.
- The model is not able to classify even with grayscale image which has reduced dimensions.
- Heavy validation losses also shows that the model struggles to distinguish between people With Helmets and those Without Helmets, as this indicates the model had not learned the distinguishing features effectively.

Model 2: (VGG-16 (Base))

```
vgg16_m=VGG16(weights='imagenet',include_top=False,input_shape=(200,200,3))
```

```
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5
58889256/58889256 _____ 0s 0us/step
```

```
for layers in vgg16_m.layers:
    layers.trainable=False
```

```
for layers in vgg16_m.layers:
    print(layers.name,layers.trainable)
```

```
input_layer_1 False
block1_conv1 False
block1_conv2 False
block1_pool False
block2_conv1 False
block2_conv2 False
block2_pool False
block3_conv1 False
block3_conv2 False
block3_conv3 False
block3_pool False
block4_conv1 False
block4_conv2 False
block4_conv3 False
block4_pool False
block5_conv1 False
block5_conv2 False
block5_conv3 False
block5_pool False
```

```
vgg16_m.summary()
```

```
Model: "vgg16"
```

Layer (type) Param #	Output Shape	
input_layer_1 (InputLayer) 0	(None, 200, 200, 3)	
block1_conv1 (Conv2D) 1,792	(None, 200, 200, 64)	
block1_conv2 (Conv2D) 36,928	(None, 200, 200, 64)	
block1_pool (MaxPooling2D) 0	(None, 100, 100, 64)	
block2_conv1 (Conv2D) 73,856	(None, 100, 100, 128)	
block2_conv2 (Conv2D) 147,584	(None, 100, 100, 128)	
block2_pool (MaxPooling2D) 0	(None, 50, 50, 128)	
block3_conv1 (Conv2D) 295,168	(None, 50, 50, 256)	
block3_conv2 (Conv2D) 590,080	(None, 50, 50, 256)	
block3_conv3 (Conv2D) 590,080	(None, 50, 50, 256)	
block3_pool (MaxPooling2D) 0	(None, 25, 25, 256)	

block4_conv1 (Conv2D)	(None, 25, 25, 512)	
1,180,160		
block4_conv2 (Conv2D)	(None, 25, 25, 512)	
2,359,808		
block4_conv3 (Conv2D)	(None, 25, 25, 512)	
2,359,808		
block4_pool (MaxPooling2D)	(None, 12, 12, 512)	
0		
block5_conv1 (Conv2D)	(None, 12, 12, 512)	
2,359,808		
block5_conv2 (Conv2D)	(None, 12, 12, 512)	
2,359,808		
block5_conv3 (Conv2D)	(None, 12, 12, 512)	
2,359,808		
block5_pool (MaxPooling2D)	(None, 6, 6, 512)	
0		

Total params: 14,714,688 (56.13 MB)

Trainable params: 0 (0.00 B)

Non-trainable params: 14,714,688 (56.13 MB)

cls()

m2=Sequential()

m2.add(vgg16_m)

m2.add(Flatten())

m2.add(Dense(1, activation='sigmoid'))

opti2=keras.optimizers.Adam(learning_rate=0.001)

m2.compile(optimizer=opti2,loss='binary_crossentropy',metrics=['accuracy'])


```
h2=m2.fit(xc_train_normalized,yc_train,epochs=32,validation_data=(xc_val_normalized,yc_val),batch_size=32)
```

Epoch 1/32

12/12 _____ 43s 3s/step - accuracy: 0.7551 - loss: 0.4722 - val_accuracy: 0.9841 - val_loss: 0.0369

Epoch 2/32

12/12 _____ 42s 211ms/step - accuracy: 1.0000 - loss: 0.0167 - val_accuracy: 0.9841 - val_loss: 0.0183

Epoch 3/32

12/12 _____ 3s 212ms/step - accuracy: 1.0000 - loss: 0.0043 - val_accuracy: 1.0000 - val_loss: 0.0104

Epoch 4/32

12/12 _____ 3s 273ms/step - accuracy: 1.0000 - loss: 0.0034 - val_accuracy: 1.0000 - val_loss: 0.0097

Epoch 5/32

12/12 _____ 5s 276ms/step - accuracy: 1.0000 - loss: 0.0026 - val_accuracy: 1.0000 - val_loss: 0.0104

Epoch 6/32

12/12 _____ 3s 275ms/step - accuracy: 1.0000 - loss: 0.0021 - val_accuracy: 1.0000 - val_loss: 0.0106

Epoch 7/32

12/12 _____ 3s 220ms/step - accuracy: 1.0000 - loss: 0.0019 - val_accuracy: 1.0000 - val_loss: 0.0101

Epoch 8/32

12/12 _____ 5s 222ms/step - accuracy: 1.0000 - loss: 0.0017 - val_accuracy: 1.0000 - val_loss: 0.0095

Epoch 9/32

12/12 _____ 3s 278ms/step - accuracy: 1.0000 - loss: 0.0016 - val_accuracy: 1.0000 - val_loss: 0.0092

Epoch 10/32

12/12 _____ 3s 282ms/step - accuracy: 1.0000 - loss: 0.0015 - val_accuracy: 1.0000 - val_loss: 0.0089

Epoch 11/32

12/12 _____ 3s 225ms/step - accuracy: 1.0000 - loss: 0.0014 - val_accuracy: 1.0000 - val_loss: 0.0087

Epoch 12/32

12/12 _____ 5s 226ms/step - accuracy: 1.0000 - loss: 0.0013 - val_accuracy: 1.0000 - val_loss: 0.0085

Epoch 13/32

12/12 _____ 5s 223ms/step - accuracy: 1.0000 - loss: 0.0012 - val_accuracy: 1.0000 - val_loss: 0.0083

Epoch 14/32

12/12 _____ 3s 278ms/step - accuracy: 1.0000 - loss: 0.0011 - val_accuracy: 1.0000 - val_loss: 0.0082

Epoch 15/32

12/12 _____ 5s 223ms/step - accuracy: 1.0000 - loss: 0.0010 - val_accuracy: 1.0000 - val_loss: 0.0080

Epoch 16/32

12/12 _____ 3s 221ms/step - accuracy: 1.0000 - loss:

9.4375e-04 - val_accuracy: 1.0000 - val_loss: 0.0078
Epoch 17/32
12/12 _____ 6s 277ms/step - accuracy: 1.0000 - loss:
8.8498e-04 - val_accuracy: 1.0000 - val_loss: 0.0077
Epoch 18/32
12/12 _____ 4s 219ms/step - accuracy: 1.0000 - loss:
8.3164e-04 - val_accuracy: 1.0000 - val_loss: 0.0075
Epoch 19/32
12/12 _____ 3s 217ms/step - accuracy: 1.0000 - loss:
7.8314e-04 - val_accuracy: 1.0000 - val_loss: 0.0074
Epoch 20/32
12/12 _____ 5s 216ms/step - accuracy: 1.0000 - loss:
7.3893e-04 - val_accuracy: 1.0000 - val_loss: 0.0073
Epoch 21/32
12/12 _____ 3s 218ms/step - accuracy: 1.0000 - loss:
6.9851e-04 - val_accuracy: 1.0000 - val_loss: 0.0072
Epoch 22/32
12/12 _____ 5s 214ms/step - accuracy: 1.0000 - loss:
6.6145e-04 - val_accuracy: 1.0000 - val_loss: 0.0070
Epoch 23/32
12/12 _____ 5s 219ms/step - accuracy: 1.0000 - loss:
6.2740e-04 - val_accuracy: 1.0000 - val_loss: 0.0069
Epoch 24/32
12/12 _____ 5s 213ms/step - accuracy: 1.0000 - loss:
5.9603e-04 - val_accuracy: 1.0000 - val_loss: 0.0068
Epoch 25/32
12/12 _____ 3s 217ms/step - accuracy: 1.0000 - loss:
5.6707e-04 - val_accuracy: 1.0000 - val_loss: 0.0068
Epoch 26/32
12/12 _____ 3s 217ms/step - accuracy: 1.0000 - loss:
5.4026e-04 - val_accuracy: 1.0000 - val_loss: 0.0067
Epoch 27/32
12/12 _____ 6s 275ms/step - accuracy: 1.0000 - loss:
5.1540e-04 - val_accuracy: 1.0000 - val_loss: 0.0066
Epoch 28/32
12/12 _____ 3s 275ms/step - accuracy: 1.0000 - loss:
4.9231e-04 - val_accuracy: 1.0000 - val_loss: 0.0065
Epoch 29/32
12/12 _____ 5s 279ms/step - accuracy: 1.0000 - loss:
4.7080e-04 - val_accuracy: 1.0000 - val_loss: 0.0064
Epoch 30/32
12/12 _____ 4s 219ms/step - accuracy: 1.0000 - loss:
4.5075e-04 - val_accuracy: 1.0000 - val_loss: 0.0063
Epoch 31/32
12/12 _____ 5s 219ms/step - accuracy: 1.0000 - loss:
4.3202e-04 - val_accuracy: 1.0000 - val_loss: 0.0063
Epoch 32/32
12/12 _____ 5s 221ms/step - accuracy: 1.0000 - loss:
4.1448e-04 - val_accuracy: 1.0000 - val_loss: 0.0062


```
m2tr=model_performance_classification(m2,xc_train_normalized,yc_train)
m2tr
```

12/12 ————— 3s 204ms/step

```
{"summary":{"\n  \"name\": \"m2tr\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"F1 Score\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  },\n  \"type\": \"dataframe\",\n  \"variable_name\": \"m2tr\"}
```

```
plot_confusion_matrix(m2,xc_train_normalized,yc_train)
```

12/12 ————— 2s 161ms/step



```
m2vl=model_performance_classification(m2,xc_val_normalized,yc_val)
m2vl
```

4/4 ————— 1s 267ms/step

```
{
  "summary": {
    "name": "m2vl",
    "rows": 1,
    "fields": [
      {
        "column": "Accuracy",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [
            1.0
          ],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "Recall",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [
            1.0
          ],
          "semantic_type": ""
        },
        "description": ""
      }
    ]
  }
}
```


```

{"Precision": 1.0, "properties": {"dtype": "number", "std": null, "min": 1.0, "max": 1.0, "num_unique_values": 1, "samples": [1.0]}, "description": "F1 Score", "column": "F1 Score", "properties": {"dtype": "number", "std": null, "min": 1.0, "max": 1.0, "num_unique_values": 1, "samples": [1.0]}, "semantic_type": "F1 Score", "description": "F1 Score"}
{"type": "dataframe", "variable_name": "m2vl"}


```

plot_confusion_matrix(m2,xc_val_normalized,yc_val)

4/4 ————— 1s 160ms/step



```
plot_history(h2)
```



Visualizing the prediction:


```
prev=m2.predict(xc_train_normalized)
f,ax=plt.subplots(4,2,figsize=(15,15))
for i,j in zip(ax.flatten(),range(8)):
    i.imshow(xc_train_normalized[j])
    i.axis('off')
    if prev[j]>0.5:
        i.set_title(f'With Helmet:{prev[j]}')
    else:
        i.set_title(f'Without Helmet:{prev[j]}')
```

12/12 ————— 2s 162ms/step

With Helmet:[0.99999464]




With Helmet:[0.997755]



Without Helmet:[2.100612e-05]




With Helmet:[0.99984515]



With Helmet:[0.9999515]




With Helmet:[0.99999917]



With Helmet:[0.99998057]



Without Helmet:[0.00033975]



Observations on model 2:

- As we see the metrics(recall) and confusion matrix the score are too good which can also be overfitting.
- while The validation loss is slightly higher than training it is within limits as the curves are smooth and uniform.
- But, the Prediction visualisations are also correct maybe due to the data being small & clean.

Model 3: (VGG-16 (Base + FFNN))

```
cls()

m3=Sequential()
m3.add(vgg16_m)
m3.add(Flatten())
m3.add(Dense(32, activation='relu'))
m3.add(BatchNormalization())
m3.add(Dropout(0.5))
m3.add(Dense(32, activation='relu'))
m3.add(BatchNormalization())
m3.add(Dropout(0.5))
m3.add(Dense(1, activation='sigmoid'))

opti=keras.optimizers.Adam(learning_rate=0.001)
m3.compile(optimizer=opti,loss='binary_crossentropy',metrics=['accuracy'])

h3=m3.fit(xc_train_normalized,yc_train,epochs=32,validation_data=(xc_val_normalized,yc_val),batch_size=32)
```

Epoch 1/32
12/12 _____ 11s 605ms/step - accuracy: 0.7833 - loss: 0.5235 - val_accuracy: 0.7222 - val_loss: 0.4812

Epoch 2/32
12/12 _____ 3s 276ms/step - accuracy: 0.9462 - loss: 0.1645 - val_accuracy: 0.8492 - val_loss: 0.2329

Epoch 3/32
12/12 _____ 5s 276ms/step - accuracy: 0.9771 - loss: 0.1006 - val_accuracy: 0.9524 - val_loss: 0.1292

Epoch 4/32
12/12 _____ 3s 279ms/step - accuracy: 0.9630 - loss: 0.0867 - val_accuracy: 1.0000 - val_loss: 0.0666

Epoch 5/32
12/12 _____ 5s 276ms/step - accuracy: 0.9935 - loss: 0.0552 - val_accuracy: 1.0000 - val_loss: 0.0408

Epoch 6/32
12/12 _____ 5s 221ms/step - accuracy: 0.9755 - loss: 0.0696 - val_accuracy: 1.0000 - val_loss: 0.0267

Epoch 7/32
12/12 _____ 5s 220ms/step - accuracy: 0.9967 - loss: 0.0378 - val_accuracy: 1.0000 - val_loss: 0.0167

Epoch 8/32
12/12 _____ 6s 275ms/step - accuracy: 0.9996 - loss: 0.0352 - val_accuracy: 1.0000 - val_loss: 0.0132

Epoch 9/32
12/12 _____ 5s 281ms/step - accuracy: 0.9917 - loss: 0.0371 - val_accuracy: 1.0000 - val_loss: 0.0098

Epoch 10/32
12/12 _____ 3s 220ms/step - accuracy: 0.9991 - loss: 0.0277 - val_accuracy: 1.0000 - val_loss: 0.0083

Epoch 11/32
12/12 _____ 5s 219ms/step - accuracy: 0.9909 - loss: 0.0320 - val_accuracy: 1.0000 - val_loss: 0.0082

Epoch 12/32
12/12 _____ 5s 224ms/step - accuracy: 1.0000 - loss: 0.0246 - val_accuracy: 0.9921 - val_loss: 0.0092

Epoch 13/32
12/12 _____ 3s 222ms/step - accuracy: 1.0000 - loss: 0.0171 - val_accuracy: 0.9921 - val_loss: 0.0093

Epoch 14/32
12/12 _____ 5s 220ms/step - accuracy: 1.0000 - loss: 0.0183 - val_accuracy: 0.9921 - val_loss: 0.0091

Epoch 15/32
12/12 _____ 3s 223ms/step - accuracy: 1.0000 - loss: 0.0196 - val_accuracy: 0.9921 - val_loss: 0.0095

Epoch 16/32
12/12 _____ 5s 221ms/step - accuracy: 0.9923 - loss: 0.0234 - val_accuracy: 0.9921 - val_loss: 0.0097

Epoch 17/32
12/12 _____ 3s 278ms/step - accuracy: 1.0000 - loss: 0.0119 - val_accuracy: 0.9921 - val_loss: 0.0115

Epoch 18/32
12/12 _____ 3s 222ms/step - accuracy: 1.0000 - loss: 0.0144 - val_accuracy: 0.9921 - val_loss: 0.0123

Epoch 19/32
12/12 _____ 3s 225ms/step - accuracy: 1.0000 - loss: 0.0123 - val_accuracy: 0.9921 - val_loss: 0.0120

Epoch 20/32
12/12 _____ 5s 220ms/step - accuracy: 1.0000 - loss: 0.0087 - val_accuracy: 0.9921 - val_loss: 0.0112

Epoch 21/32
12/12 _____ 6s 278ms/step - accuracy: 1.0000 - loss: 0.0079 - val_accuracy: 0.9921 - val_loss: 0.0112

Epoch 22/32
12/12 _____ 3s 278ms/step - accuracy: 0.9989 - loss: 0.0151 - val_accuracy: 0.9921 - val_loss: 0.0109

Epoch 23/32

12/12 _____ 4s 218ms/step - accuracy: 1.0000 - loss: 0.0099 - val_accuracy: 0.9921 - val_loss: 0.0117

Epoch 24/32

12/12 _____ 3s 220ms/step - accuracy: 1.0000 - loss: 0.0083 - val_accuracy: 0.9921 - val_loss: 0.0128

Epoch 25/32

12/12 _____ 3s 221ms/step - accuracy: 1.0000 - loss: 0.0087 - val_accuracy: 0.9921 - val_loss: 0.0130

Epoch 26/32

12/12 _____ 5s 218ms/step - accuracy: 1.0000 - loss: 0.0093 - val_accuracy: 0.9921 - val_loss: 0.0129

Epoch 27/32

12/12 _____ 3s 220ms/step - accuracy: 0.9994 - loss: 0.0077 - val_accuracy: 0.9921 - val_loss: 0.0126

Epoch 28/32

12/12 _____ 6s 279ms/step - accuracy: 0.9994 - loss: 0.0110 - val_accuracy: 0.9921 - val_loss: 0.0138

Epoch 29/32

12/12 _____ 3s 219ms/step - accuracy: 1.0000 - loss: 0.0115 - val_accuracy: 0.9921 - val_loss: 0.0141

Epoch 30/32

12/12 _____ 5s 218ms/step - accuracy: 1.0000 - loss: 0.0074 - val_accuracy: 0.9921 - val_loss: 0.0157

Epoch 31/32

12/12 _____ 3s 220ms/step - accuracy: 1.0000 - loss: 0.0142 - val_accuracy: 0.9921 - val_loss: 0.0175

Epoch 32/32

12/12 _____ 5s 219ms/step - accuracy: 1.0000 - loss: 0.0124 - val_accuracy: 0.9921 - val_loss: 0.0175

m3tr=model_performance_classification(m3,xc_train_normalized,yc_train)
m3tr

12/12 _____ 3s 213ms/step


```
{
  "summary": {
    "name": "m3tr",
    "rows": 1,
    "fields": [
      {
        "column": "Accuracy",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "Recall",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "Precision",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      }
    ]
  }
}
```




```
{
  "summary": {
    "name": "m3vl",
    "rows": 1,
    "fields": [
      {
        "column": "Accuracy",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 0.9920634920634921,
          "max": 0.9920634920634921,
          "num_unique_values": 1,
          "samples": [
            0.9920634920634921
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Recall",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 0.9920634920634921,
          "max": 0.9920634920634921,
          "num_unique_values": 1,
          "samples": [
            0.9920634920634921
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "Precision",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 0.9921894683799446,
          "max": 0.9921894683799446,
          "num_unique_values": 1,
          "samples": [
            0.9921894683799446
          ],
          "semantic_type": "",
          "description": ""
        },
        "column": "F1 Score",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 0.9920639920009998,
          "max": 0.9920639920009998,
          "num_unique_values": 1,
          "samples": [
            0.9920639920009998
          ],
          "semantic_type": "",
          "description": ""
        }
      ]
    }
  },
  "type": "dataframe",
  "variable_name": "m3vl"
}
```

```
plot_confusion_matrix(m3, xc_val_normalized, yc_val)
```

4/4 ————— 1s 162ms/step



```
plot_history(h3)
```



Visualizing the predictions


```
prev=m3.predict(xc_train_normalized)
f,ax=plt.subplots(4,2,figsize=(15,15))
for i,j in zip(ax.flatten(),range(8)):
    i.imshow(xc_train_normalized[j])
    i.axis('off')
    if prev[j]>0.5:
        i.set_title(f'With Helmet:{prev[j]}')
    else:
        i.set_title(f'Without Helmet:{prev[j]}')
```

12/12 ————— 2s 158ms/step

With Helmet:[0.99993825]



With Helmet:[0.9992337]



Without Helmet:[0.00010013]




With Helmet:[0.99967813]



With Helmet:[0.9998204]




With Helmet:[0.9999492]



With Helmet:[0.99994075]



Without Helmet:[0.00035484]



Observations of model 3::

- The model also shows slight overfitting.
- The loss curves show slight oscillations when compared to **Model 2**, but there is convergence and the validation losses are smaller in this model.
- The Predicted values are becoming more concise in certain cases with metric evaluation being same as previous model.
- As the metrics(recall) and confusion matrix scores show model is doing really well may not be robust since the data being clean and small.

Model 4: (VGG-16 (Base + FFNN + Data Augmentation))

- In most of the real-world case studies, it is challenging to acquire a large number of images and then train CNNs.
- To overcome this problem, one approach we might consider is **Data Augmentation**.
- CNNs have the property of **translational invariance**, which means they can recognise an object even if its appearance shifts translationally in some way. - Taking this attribute into account, we can augment the images using the techniques listed below
 - Horizontal Flip (should be set to True/False)
 - Vertical Flip (should be set to True/False)
 - Height Shift (should be between 0 and 1)
 - Width Shift (should be between 0 and 1)
 - Rotation (should be between 0 and 180)
 - Shear (should be between 0 and 1)
 - Zoom (should be between 0 and 1) etc.

Remember, **data augmentation should not be used in the validation/test data set.**

```
train_datagen = ImageDataGenerator(
    horizontal_flip = True,
    vertical_flip = False,
    height_shift_range= 0.1,
    width_shift_range=0.1,
    rotation_range=20,
    shear_range = 0.1,
    zoom_range=0.1)

# Flowing training images in batches of 20 using train_datagen
# generator
train_generator =
train_datagen.flow(xc_train_normalized,yc_train,batch_size=21,shuffle=
True)

cls()
```



```

m4=Sequential()
m4.add(vgg16_m)
m4.add(Flatten())
m4.add(Dense(32, activation='relu'))
m4.add(BatchNormalization())
m4.add(Dropout(0.5))
m4.add(Dense(32, activation='relu'))
m4.add(BatchNormalization())
m4.add(Dropout(0.5))
m4.add(Dense(1, activation='sigmoid'))

opti=keras.optimizers.Adam(learning_rate=0.001)
m4.compile(optimizer=opti,loss='binary_crossentropy',metrics=['accuracy'])

h4=m4.fit(train_generator,epochs=32,validation_data=(xc_val_normalized,yc_val),batch_size=32)

Epoch 1/32
18/18 _____ 19s 387ms/step - accuracy: 0.8195 - loss: 0.3973 - val_accuracy: 1.0000 - val_loss: 0.0460
Epoch 2/32
18/18 _____ 5s 288ms/step - accuracy: 0.9513 - loss: 0.1421 - val_accuracy: 1.0000 - val_loss: 0.0367
Epoch 3/32
18/18 _____ 5s 248ms/step - accuracy: 0.9645 - loss: 0.1034 - val_accuracy: 0.9921 - val_loss: 0.0302
Epoch 4/32
18/18 _____ 6s 320ms/step - accuracy: 0.9726 - loss: 0.0863 - val_accuracy: 0.9921 - val_loss: 0.0200
Epoch 5/32
18/18 _____ 5s 251ms/step - accuracy: 0.9960 - loss: 0.0553 - val_accuracy: 0.9921 - val_loss: 0.0168
Epoch 6/32
18/18 _____ 5s 290ms/step - accuracy: 0.9753 - loss: 0.0517 - val_accuracy: 0.9921 - val_loss: 0.0203
Epoch 7/32
18/18 _____ 6s 315ms/step - accuracy: 0.9937 - loss: 0.0350 - val_accuracy: 0.9841 - val_loss: 0.0285
Epoch 8/32
18/18 _____ 5s 287ms/step - accuracy: 0.9948 - loss: 0.0327 - val_accuracy: 0.9921 - val_loss: 0.0224
Epoch 9/32
18/18 _____ 5s 279ms/step - accuracy: 0.9928 - loss: 0.0441 - val_accuracy: 0.9921 - val_loss: 0.0211
Epoch 10/32
18/18 _____ 5s 258ms/step - accuracy: 0.9993 - loss: 0.0189 - val_accuracy: 0.9921 - val_loss: 0.0160
Epoch 11/32
18/18 _____ 5s 250ms/step - accuracy: 0.9879 - loss:

```

0.0284 - val_accuracy: 0.9921 - val_loss: 0.0096
Epoch 12/32
18/18 _____ 6s 307ms/step - accuracy: 0.9936 - loss:
0.0318 - val_accuracy: 0.9921 - val_loss: 0.0123
Epoch 13/32
18/18 _____ 5s 250ms/step - accuracy: 0.9946 - loss:
0.0231 - val_accuracy: 0.9921 - val_loss: 0.0138
Epoch 14/32
18/18 _____ 5s 285ms/step - accuracy: 1.0000 - loss:
0.0116 - val_accuracy: 0.9921 - val_loss: 0.0122
Epoch 15/32
18/18 _____ 10s 248ms/step - accuracy: 0.9937 - loss:
0.0388 - val_accuracy: 0.9921 - val_loss: 0.0110
Epoch 16/32
18/18 _____ 6s 332ms/step - accuracy: 1.0000 - loss:
0.0148 - val_accuracy: 0.9921 - val_loss: 0.0117
Epoch 17/32
18/18 _____ 5s 248ms/step - accuracy: 0.9979 - loss:
0.0221 - val_accuracy: 0.9921 - val_loss: 0.0132
Epoch 18/32
18/18 _____ 5s 251ms/step - accuracy: 0.9917 - loss:
0.0255 - val_accuracy: 0.9921 - val_loss: 0.0109
Epoch 19/32
18/18 _____ 6s 312ms/step - accuracy: 0.9976 - loss:
0.0119 - val_accuracy: 0.9921 - val_loss: 0.0091
Epoch 20/32
18/18 _____ 4s 247ms/step - accuracy: 0.9984 - loss:
0.0122 - val_accuracy: 0.9921 - val_loss: 0.0099
Epoch 21/32
18/18 _____ 5s 287ms/step - accuracy: 0.9979 - loss:
0.0178 - val_accuracy: 0.9921 - val_loss: 0.0113
Epoch 22/32
18/18 _____ 5s 271ms/step - accuracy: 0.9887 - loss:
0.0341 - val_accuracy: 0.9921 - val_loss: 0.0191
Epoch 23/32
18/18 _____ 5s 250ms/step - accuracy: 1.0000 - loss:
0.0113 - val_accuracy: 0.9921 - val_loss: 0.0199
Epoch 24/32
18/18 _____ 5s 304ms/step - accuracy: 1.0000 - loss:
0.0123 - val_accuracy: 0.9921 - val_loss: 0.0134
Epoch 25/32
18/18 _____ 5s 251ms/step - accuracy: 0.9957 - loss:
0.0289 - val_accuracy: 0.9921 - val_loss: 0.0122
Epoch 26/32
18/18 _____ 5s 287ms/step - accuracy: 0.9972 - loss:
0.0244 - val_accuracy: 1.0000 - val_loss: 0.0025
Epoch 27/32
18/18 _____ 6s 308ms/step - accuracy: 0.9920 - loss:
0.0244 - val_accuracy: 1.0000 - val_loss: 0.0018

```

Epoch 28/32
18/18 _____ 5s 250ms/step - accuracy: 0.9997 - loss:
0.0184 - val_accuracy: 1.0000 - val_loss: 0.0019
Epoch 29/32
18/18 _____ 5s 257ms/step - accuracy: 1.0000 - loss:
0.0089 - val_accuracy: 1.0000 - val_loss: 0.0030
Epoch 30/32
18/18 _____ 5s 281ms/step - accuracy: 0.9993 - loss:
0.0083 - val_accuracy: 1.0000 - val_loss: 0.0053
Epoch 31/32
18/18 _____ 5s 287ms/step - accuracy: 0.9997 - loss:
0.0135 - val_accuracy: 1.0000 - val_loss: 0.0046
Epoch 32/32
18/18 _____ 6s 307ms/step - accuracy: 1.0000 - loss:
0.0048 - val_accuracy: 0.9921 - val_loss: 0.0066

```

```
# Initialize empty lists to collect data
```

```
all_images = []
all_labels = []
```

```
# Iterate through the generator until all data is retrieved
```

```
for _ in range(len(train_generator)):
    imag, labela = next(train_generator)
    all_images.append(imag)
    all_labels.append(labela)
```

```
# Concatenate all batches into single arrays
```

```
all_images = np.concatenate(all_images, axis=0)
all_labels = np.concatenate(all_labels, axis=0)
```

```
print(all_images.shape) # Shape: (total_samples, height, width,
channels)
```

```
print(all_labels.shape) #
```

```
(378, 200, 200, 3)
(378,)
```

```
all_labels=pd.Series(all_labels)
all_labels
```

```

0      1
1      0
2      1
3      0
4      1
..
373    0
374    1
375    1
376    1

```

```
377      0
Length: 378, dtype: int64
```


```
m4tr=model_performance_classification(m4,all_images,all_labels)
m4tr
```

```
12/12 ————— 3s 205ms/step
```

```
{
  "summary": {
    "name": "m4tr",
    "rows": 1,
    "fields": [
      {
        "column": "Accuracy",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "Recall",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "Precision",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "F1 Score",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 1.0,
          "max": 1.0,
          "num_unique_values": 1,
          "samples": [1.0],
          "semantic_type": ""
        },
        "description": ""
      }
    ]
  },
  "type": "dataframe",
  "variable_name": "m4tr"
}
```

```
plot_confusion_matrix(m4,all_images,all_labels)
```


```
12/12 ————— 2s 157ms/step
```




```
m4vl=model_performance_classification(m4,xc_val_normalized,yc_val)
m4vl
```

4/4 ————— 1s 307ms/step

```
{
  "summary": {
    "name": "m4vl",
    "rows": 1,
    "fields": [
      {
        "column": "Accuracy",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 0.9920634920634921,
          "max": 0.9920634920634921,
          "num_unique_values": 1,
          "samples": [
            0.9920634920634921
          ],
          "semantic_type": ""
        },
        "description": ""
      },
      {
        "column": "Recall",
        "properties": {
          "dtype": "number",
          "std": null,
          "min": 0.9920634920634921,
          "max": 0.9920634920634921,
          "num_unique_values": 1,
          "samples": [
            0.9920634920634921
          ]
        },
        "description": ""
      }
    ]
  }
}
```

```
plot_history(h4)
```



Visualizing the predictions

```
prev=m4.predict(imag)

f,ax=plt.subplots(4,2,figsize=(15,15))
for i,j in zip(ax.flatten(),range(8)):
    i.imshow(imag[j])
    i.axis('off')
    if prev[j]>0.5:
        i.set_title(f'With Helmet:{prev[j]}')
    else:
        i.set_title(f'Without Helmet:{prev[j]}')

1/1 ————— 1s 516ms/step
```


With Helmet:[0.99971765]



With Helmet:[0.99994946]




With Helmet:[0.9992712]




Without Helmet:[5.3697964e-05]



With Helmet:[0.9997824]



Without Helmet:[0.00016988]



With Helmet:[0.99954873]



With Helmet:[0.99994206]



Observations on model 4:

- Although model may shows slight overfitting The metrics(recall) and matix shows the model is doing very well even with the augmented data which shows its robustness when compared to the previous models.
- The loss curves suggest that validation losses are more smaller than previous models.

Model Performance Comparison and Final Model Selection

```
train_eval_result=pd.concat([m1tr.T,m2tr.T,m3tr.T,m4tr.T],axis=1)
train_eval_result.columns=['Model 1:Basic CNN','Model 2:Base
VGG16','Model 3:VGG16+FNN','Model 4:VGG16+FNN+Data augmentation']
print("Models Results on Training data")
print("*****50)
train_eval_result
```

Models Results on Training data

```
*****
*****
```

```
{
  "summary": {
    "\n  \"name\": \"train_eval_result\",
    "\n  \"rows\": 4,
    "\n  \"fields\": [
      {\n        \"column\": \"Model 1:Basic CNN\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 0.11096318684183888,
          \"min\": 0.5115102748836582,
          \"max\": 0.7747412008281573,
          \"num_unique_values\": 3,
          \"samples\": [
            0.5952380952380952,
            0.7747412008281573,
            0.5115102748836582
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      },
      {\n        \"column\": \"Model 2:Base VGG16\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 0.0,
          \"min\": 1.0,
          \"max\": 1.0,
          \"num_unique_values\": 1,
          \"samples\": [
            1.0
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      },
      {\n        \"column\": \"Model 3:VGG16+FNN\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 0.0,
          \"min\": 1.0,
          \"max\": 1.0,
          \"num_unique_values\": 1,
          \"samples\": [
            1.0
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      },
      {\n        \"column\": \"Model 4:VGG16+FNN+Data augmentation\",
        \"properties\": {
          \"dtype\": \"number\",
          \"std\": 0.0,
          \"min\": 1.0,
          \"max\": 1.0,
          \"num_unique_values\": 1,
          \"samples\": [
            1.0
          ],
          \"semantic_type\": \"\",
          \"description\": \"\"
        }
      }
    ]
  },
  \"type\": \"dataframe\",
  \"variable_name\": \"train_eval_result\"
}
```

```

val_eval_result=pd.concat([m1vl.T,m2vl.T,m3vl.T,m4vl.T],axis=1)
val_eval_result.columns=train_eval_result.columns
print("Models Results on validation data")
print("***50)
val_eval_result

```

Models Results on validation data

```

*****
*****

```

```

{"summary":{"\n  \"name\": \"val_eval_result\",\n  \"rows\": 4,\n  \"fields\": [\n    {\n      \"column\": \"Model 1:Basic CNN\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.10250210310405837,\n        \"min\": 0.5379033270558695,\n        \"max\": 0.7797443461160275,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.6111111111111112,\n          0.7797443461160275,\n          0.5379033270558695\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Model 2:Base VGG16\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.0,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Model 3:VGG16+FNN\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 6.290527678051543e-05,\n        \"min\": 0.9920634920634921,\n        \"max\": 0.9921894683799446,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.9920634920634921\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    },\n    {\n      \"column\": \"Model 4:VGG16+FNN+Data augmentation\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 6.290527678051543e-05,\n        \"min\": 0.9920634920634921,\n        \"max\": 0.9921894683799446,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.9920634920634921\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    }\n  ],\n  \"type\":\"dataframe\", \"variable_name\":\"val_eval_result\"}

```

Final Model Selection

Based on the observations:

- As we observe model 1 metrics are not good and model is not able to classify people with helmets.
- The results of **models 2, 3, and 4** are identical in terms of performance metrics(recall).
- However, **Model 4** was trained using augmented data, which enhances its robustness to variations in the images.

Conclusion

Model 4 is selected as the final model for deployment due to its ability to generalize better to unseen data.

Test Performance

```
m4ts=model_performance_classification(m4,xc_test_normalized,yc_test)
m4ts
```

4/4 ————— 10s 3s/step

```
{"summary":{"\n  \"name\": \"m4ts\",\n  \"rows\": 1,\n  \"fields\": [\n    {\n      \"column\": \"Accuracy\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Recall\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Precision\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"F1 Score\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": null,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  },\n  \"type\": \"dataframe\",\n  \"variable_name\": \"m4ts\"}
```


```
m4eval=pd.concat([m4tr.T,m4vl.T,m4ts.T],axis=1)
m4eval.columns=['Train','Validation','Test']
m4eval
```

```
{"summary":{"\n  \"name\": \"m4eval\",\n  \"rows\": 4,\n  \"fields\": [\n    {\n      \"column\": \"Train\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 0.0,\n        \"min\": 1.0,\n        \"max\": 1.0,\n        \"num_unique_values\": 1,\n        \"samples\": [\n          1.0\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Validation\",\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 6.290527678051543e-05,\n        \"min\": 0.9920634920634921,\n        \"max\": 0.9921894683799446,\n        \"num_unique_values\": 3,\n        \"samples\": [\n          0.9920634920634921\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": \"Test\",\n      \"properties\": {\n        \"dtype\": \"number\",
```


```
\ "std\ ": 0.0,\n      \ "min\ ": 1.0,\n      \ "max\ ": 1.0,\n \ "num_unique_values\ ": 1,\n      \ "samples\ ": [\n      1.0\n ],\n      \ "semantic_type\ ": \ "\",\n      \ "description\ ": \ "\">\n }\n   }\n   ]\n }", "type": "dataframe", "variable_name": "m4eval" }
```

```
plot_confusion_matrix(m4,xc_test_normalized,yc_test)
```

4/4 ————— 1s 163ms/step



```
plot_history(h4)
```



```
print("Classification Report - Train data",end="\n\n")
crtr =
classification_report(yc_train,m4.predict(xc_train_normalized)>0.5)
print(crtr)
```

Classification Report - Train data

```
12/12 ----- 2s 160ms/step
              precision    recall  f1-score   support

    0           1.00        1.00        1.00        192
    1           1.00        1.00        1.00        186

 accuracy          1.00          1.00          1.00        378
 macro avg          1.00          1.00          1.00        378
weighted avg          1.00          1.00          1.00        378
```

```
print("Classification Report - Validation data",end="\n\n")
crvl = classification_report(yc_val,m4.predict(xc_val_normalized)>0.5)
print(crvl)
```

Classification Report - Validation data

```
4/4 ----- 1s 159ms/step
              precision    recall  f1-score   support

    0           1.00        0.98        0.99        64
    1           0.98        1.00        0.99        62

 accuracy          0.99          0.99          0.99        126
 macro avg          0.99          0.99          0.99        126
weighted avg          0.99          0.99          0.99        126
```

```
print("Classification Report - Test data",end="\n\n")
crts =
classification_report(y_test,m4.predict(x_test_normalized)>0.5)
print(crts)
```

Classification Report - Test data

4/4	<hr/>				1s 159ms/step
		precision	recall	f1-score	support
0	1.00	1.00	1.00	1.00	64
1	1.00	1.00	1.00	1.00	63
accuracy				1.00	127
macro avg	1.00	1.00	1.00	1.00	127
weighted avg	1.00	1.00	1.00	1.00	127

Observations on best model(model 4):

- The reports, metrics(recall) and matrix as well loss curves all suggest very good performance.
- Since the model is trained on augmented data it is robust and able learn the right elements excluding the noise.

Actionable Insights & Recommendations

1. Model Evaluation and Selection

- Even though the primary evaluation criterion is recall due to the safety-critical nature of the task, all metrics (accuracy, precision, and F1-score) are essential for a holistic assessment. Model 4 achieves very good performance across these metrics with slight overfitting, making it the most suitable candidate.
- Validation results indicate that Model 4 is more robust than other models, showing convergence in loss curves and better generalization even with data augmentation which confirms robustness.

2. Data Considerations

- The dataset size is a significant limitation. Even with augmentation, it is insufficient for robust generalization. Increasing the dataset size with real-world samples is critical.
- A larger and more diverse dataset will improve the model's ability to handle varied conditions such as different lighting, backgrounds, and helmet types.
- Focus on collecting edge-case scenarios (e.g., partially visible helmets, different helmet styles, and obstructions) to ensure comprehensive training and evaluation.

3. Leveraging Advanced Architectures and Strategies

- Utilize advanced architectures, such as transformers, which excel at capturing global context and intricate patterns in data.
- Leverage models designed for improved feature extraction and adaptability, particularly beneficial for datasets with limited diversity.
- Integrate these architectures with effective preprocessing and augmentation techniques to maximize performance and reduce overfitting.

